SFL Scientific

What is a Data Lake?

- What are the benefits?
- How does a Data Lake differ from a Data Warehouse?
- How are they related?
- What are the benefits?

Data Lakes offer a flexible, scalable, and cost effective data integration solution

A **Data Lake** is a **centralized** and **scalable** repository which serves as a **landing zone** for all **incoming data streams**.

Data Lakes enable **persistent storage** of **raw data** in any format such as **unstructured** (images, video, audio), **semi-structured** (XML, JSON, HTML) or **structured** (csvs, relational data).

Data Lakes are **cost effective** since data is stored raw data and doesn't require processing.

Data Lake providers:

*

Data Lake

- Raw Data
- Schema free and highly flexible
- Difficult to query
- Data Engineer & Data Scientists
- Lower cost solution

Data Warehouse

- Processed and transformed data
- Structured schemas
- Optimized for query performance
- Marketing, Finance, and Business Intelligence
- Higher cost solution

Cleaning

Processing

Business Deliverable

Serverless Architecture

- What are the benefits?
- What are the drawbacks?

Serverless architectures provide a dynamic infrastructure solution that relieves the worry of managing resources

Serverless architectures **DO NOT** refer to compute without servers.

An **infrastructure paradigm** that enables developers to build and deploy applications **without managing** underlying servers.

This enables developers to **focus on code, system design, and functions** and removes the need to consider resource scaling and infrastructure management.

Serverless architecture is a **core concept of cloud computing** and **enables resources to scale** based on the dynamic needs of your business.

Events such as database changes, file uploads, or website requests **trigger the functionality** of the underlying microservices and applications.

Benefits

- Horizontal Scale & Redundancies
- Cost-Efficient
- Reduced Overhead
- Event-Driven Design
- Dynamic and Resource Efficient
- Quick Deployments

Drawbacks

- Requires an internet connection
- 3rd Party Dependency
- Challenges with Testing and Debugging
- Learning Curve
- Cold Starts
- Vendor Lock-In

AWS ETL Pipeline

Serverless AWS ETL Pipeline

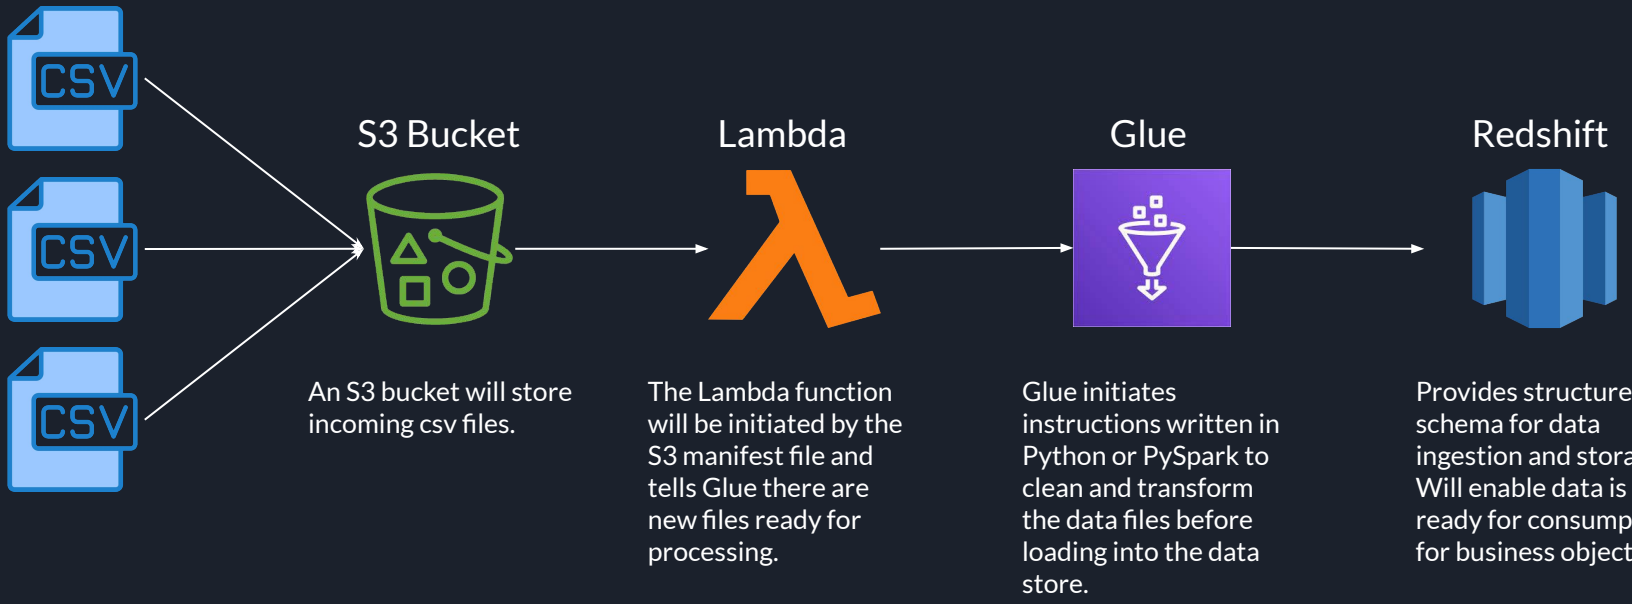

The **ETL Pipeline** designed above will receive CSV files from various sources and use an **S3 Bucket** as an incoming data store for these raw data files. Upon CSV ingestion, the **Lambda Function** will notify **AWS Glue** there is a new file ready to be processed. **Glue** will kick off a pre-defined Python recipe to transform the data so it is usable and structured. Then, it can be stored in **Redshift** where it can be used to achieve its business deliverable.

Modern MLOps

Tools and Systems for end-to-end ML product management

MLOps describes the systems and tools to manage the lifecycle of your AI & ML products

The AI & ML product lifecycle consists of four keys areas: **Development, Deployment, Monitoring, and Management**. MLOps considers the requirements at each stage.

During **development**, Data Scientists need **managed data**, **development environments,** and clearly **communicated business objectives** of what they are aiming to deliver.

Deployment of models are critical and require **controlled release environments**, **automated testing**, and continuous integration and deployment **(CI/CD)** to ensure models will perform as expected.

Once models are in production, **monitoring is paramount**. Keeping track of both **model performance** and the delivery of **business objectives** will ensure business deliverables are being met.

Model management ensures throughput can **scale** with demand, that model **version controls** are in place, and there is sufficient **documentation** for technical support and customer understanding.

MLOps Tool Suite

Communication: Tools such as Slack and GChat ensure team members can share ideas freely and stay aligned on business objectives.

Development: Jupyter & DataBricks Notebooks give Data Scientists the tools needed explore data and test theories.

Data Management: Data Lakes and data registries such as Alation ensure data assets available for use are defined and accessible.

Environment Management: Docker and Terraform are tools that can ensure consistent deployment environments so models perform as expected and minimize dependency concerns.

CI/CD: Jenkins or GitHub Actions provide build cycles for testing code changes prior to being released into production.

Monitoring: Tableau, PowerBI, and AWS Quicksight provide dashboards that can provide on demand performance regarding model metrics and business KPIs.

Orchestration: Airflow, Luigi, and AWS Glue can offer data pipeline orchestration services to manage data processes.

Scaling: Cloud providers such as AWS, Azure, and GCP provide serverless architectures to ensure your infrastructure can scale with your product's demands.

Product Evolution: MLFlow offers a model registry to ensure updated models are delivering improvements and rollbacks can be clearly determined.

Documentation: Jira and Confluence provide the tools to communicate findings and procedures, while managing tasks and deliverables of your development teams.